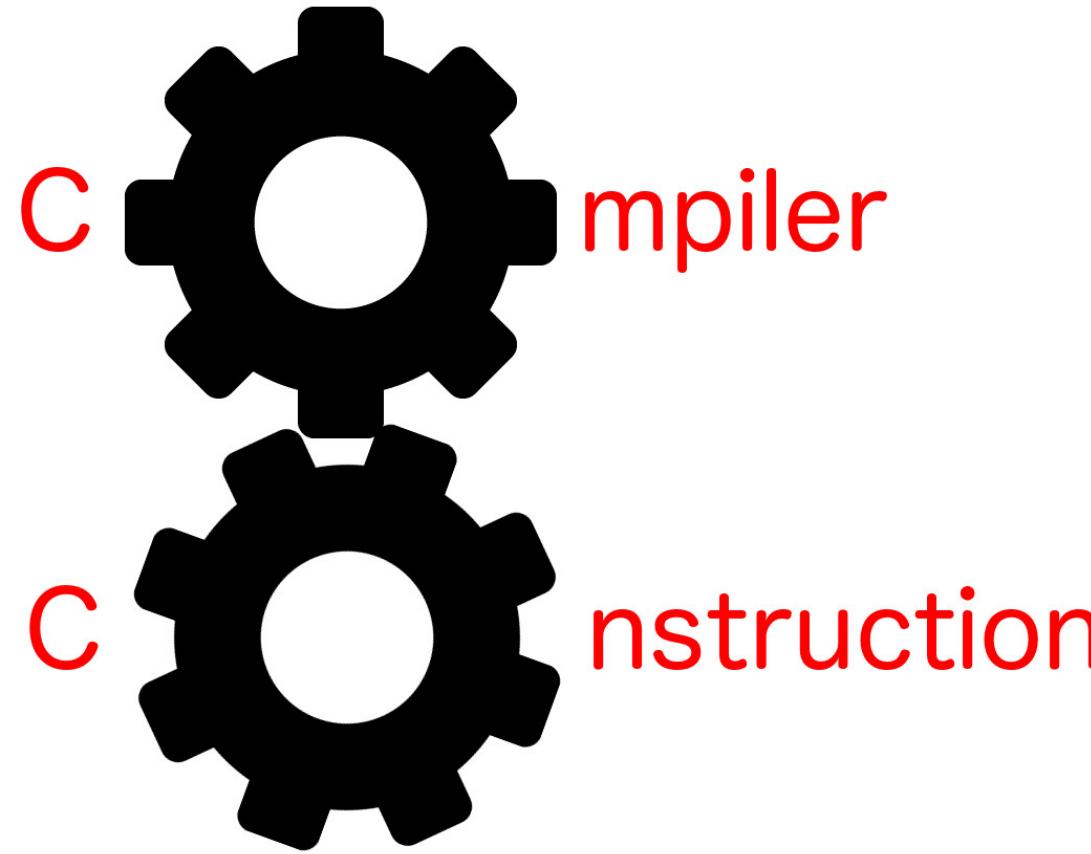




Simone Campanoni  
[simonec@eecs.northwestern.edu](mailto:simonec@eecs.northwestern.edu)



L1 -> x86\_64



# Before we start

- We use AT&T assembly syntax
  - For compatibility with GNU tools
- `rdi += rsi`
  - AT&T: `addq %rsi, %rdi`
  - Intel: `addq %rdi, %rsi`

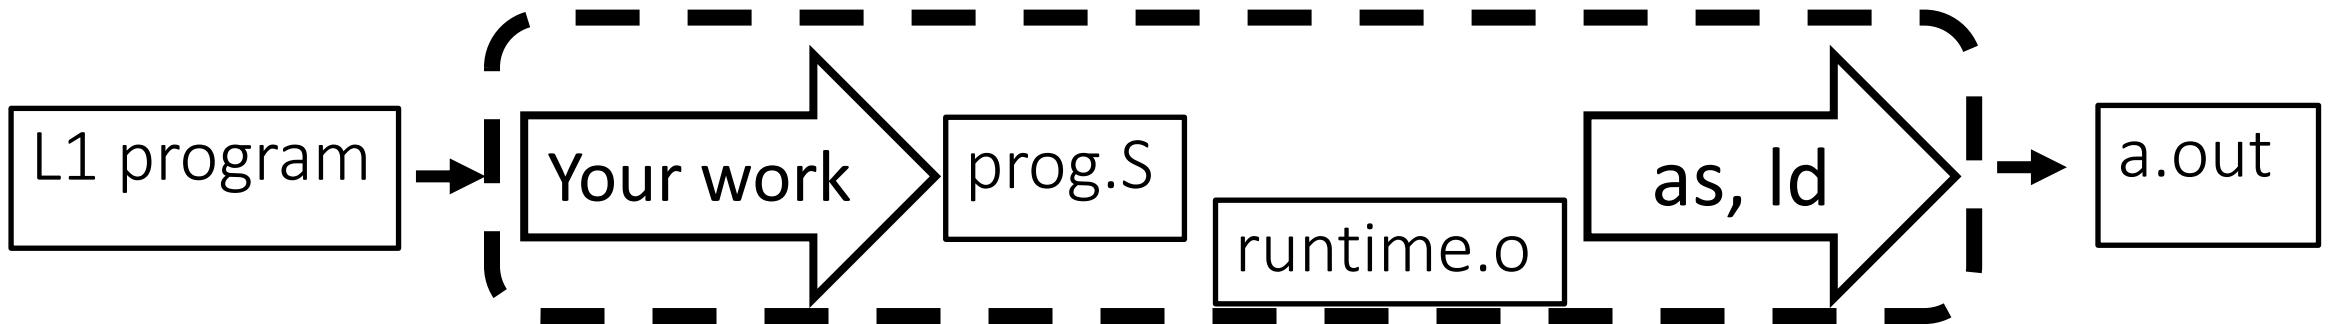
# Outline

- Setup
- From L1 to x86\_64
- Calling convention

L1c  
Makefile  
bin  
src  
tests

# Setup

- You have the structure of a compiler to start from
- Write your assignment in C++ and store the files in “src”
- You work: save x86\_64 instructions in prog.S
- The script “L1c” invokes the assembler and the linker to generate an executable binary a.out from prog.S



# A simple (incomplete) example

- Write src/compiler.cpp

```
int main(  
    int argc,  
    char **argv  
{  
    std::ofstream outputFile;  
    outputFile.open("prog.S");  
    outputFile << ".text\n" ;  
    outputFile.close();  
    return 0;  
}
```

# prog.S structure

```
.text
.globl go

go:
    # save callee-saved registers
    pushq  %rbx
    pushq  %rbp
    pushq  %r12
    pushq  %r13
    pushq  %r14
    pushq  %r15

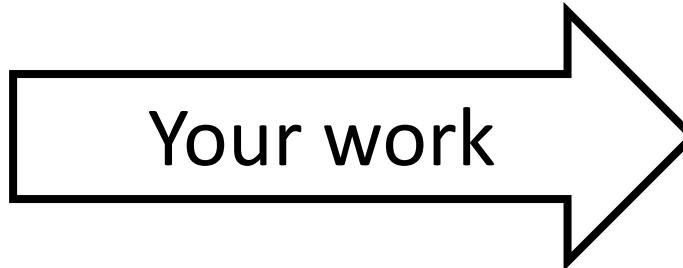
    call  «main label»

    # restore callee-saved registers and return
    popq  %r15
    popq  %r14
    popq  %r13
    popq  %r12
    popq  %rbp
    popq  %rbx
    retq
```

- runtime.c has main
- runtime.c invokes go()

# Example of prog.S

```
(:myGo  
(:myGo  
0 0  
return  
)  
)
```



```
.text  
.globl go  
go:  
    pushq %rbx  
    pushq %rbp  
    pushq %r12  
    pushq %r13  
    pushq %r14  
    pushq %r15  
    call _myGo  
    popq %r15  
    popq %r14  
    popq %r13  
    popq %r12  
    popq %rbp  
    popq %rbx  
    retq  
_myGo:  
    retq
```

The assembly code is annotated with red arrows pointing to specific instructions:

- A red arrow points to the `call _myGo` instruction.
- A red arrow points to the `retq` instruction at the end of the `go:` block.
- A red arrow points to the `retq` instruction at the end of the `_myGo:` block.

# L1

p ::= (label f<sup>+</sup>)  
f ::= (label N N i<sup>+</sup>)  
i ::= w <- s | w <- mem x M | mem x M <- s |  
     w aop t | w sop sx | w sop N | mem x M += t | mem x M -= t | w += mem x M | w -= mem x M |  
     w <- t cmp t | cjump t cmp t label label | cjump t cmp t label | label | goto label |  
     return | call u N | call print 1 | call allocate 2 | call array-error 2 |  
     w++ | w-- | w @ w w E  
w ::= a | rax | rbx | rbp | r10 | r11 | r12 | r13 | r14 | r15  
a ::= rdi | rsi | rdx | sx | r8 | r9  
sx ::= rcx  
s ::= t | label  
t ::= x | N  
u ::= w | label  
x ::= w | rsp  
aop ::= += | -= | \*= | &= |  
sop ::= <<= | >>= |  
cmp ::= < | <= | = |  
E ::= 0 | 2 | 4 | 8  
M ::= N times 8  
label ::= sequence of chars matching :[a-zA-Z\_][a-zA-Z\_0-9]\*

# Outline

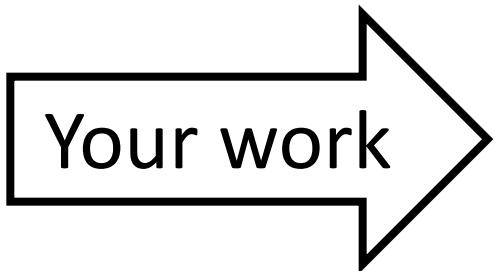
- Setup
- From L1 to x86\_64
- Calling convention

# L1 returns

To compile return instructions:

- Add **q** to specify 8 bytes values are returned

return

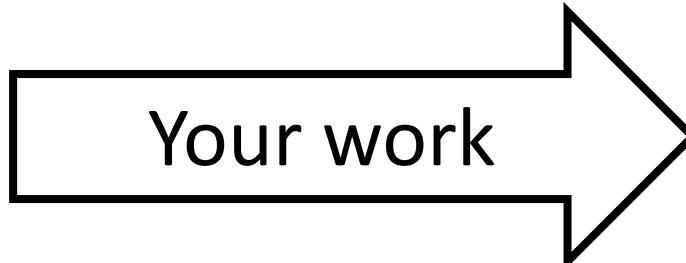


... # see later

retq

# Example of prog.S

```
(:myGo  
(:myGo  
0 0  
return  
)  
)
```



```
.text  
.globl go  
go:  
    pushq %rbx  
    pushq %rbp  
    pushq %r12  
    pushq %r13  
    pushq %r14  
    pushq %r15  
    call _myGo  
    popq %r15  
    popq %r14  
    popq %r13  
    popq %r12  
    popq %rbp  
    popq %rbx  
    retq  
_myGo:  
    retq
```

# L1 assignments

To compile simple assignments:

- prefix registers with % and constants and labels with \$
- Substitute : of labels with \_

rax <- 1

movq \$1, %rax

rax <- rbx

Your work

movq %rbx, %rax

rax <- :f

movq \$\_f, %rax

# L1 assignment example

```
(:myGo
(:myGo
0 0
rdi <- 5
return
)
)
```



```
.text
.globl go
go:
    pushq %rbx
    ...
    call _myGo
    popq %r15
    ...
    retq
_myGo:
    movq $5, %rdi ←
    retq
```

# L1 assignments to/from memory

To compile memory references:

- put parents around the register and prefix it with the offset

mem rsp 0 <- rdi

movq %rdi, 0(%rsp)

Your work

rdi <- mem rsp 8

movq 8(%rsp), %rdi

# L1 arithmetic operations

Each of the **aop=** operations correspond to their own assembly instruction

**rdi += rex**     $\Rightarrow$     **addq %rax, %rdi**

**rdi -= rex**     $\Rightarrow$     **subq %rax, %rdi**

**r10 \*= r12**     $\Rightarrow$     **imulq %r12, %r10**

**r14 &= r15**     $\Rightarrow$     **andq %r15, %r14**

# L1 arithmetic operations (2)

- `rdi--`                           $\Rightarrow$             `dec %rdi`

- `rdi++`                           $\Rightarrow$             `inc %rdi`

# L1 arithmetic operations in memory

- $\text{rdi} -= \text{mem } \text{rsp } 8$        $\Rightarrow$       `subq 8(%rsp), %rdi`
- $\text{rdi} += \text{mem } \text{rsp } 8$        $\Rightarrow$       `addq 8(%rsp), %rdi`
- $\text{mem } \text{rsp } 8 -= \text{rdi}$        $\Rightarrow$       `subq %rdi, 8(%rsp)`
- $\text{mem } \text{rsp } 8 += \text{rdi}$        $\Rightarrow$       `addq %rdi, 8(%rsp)`

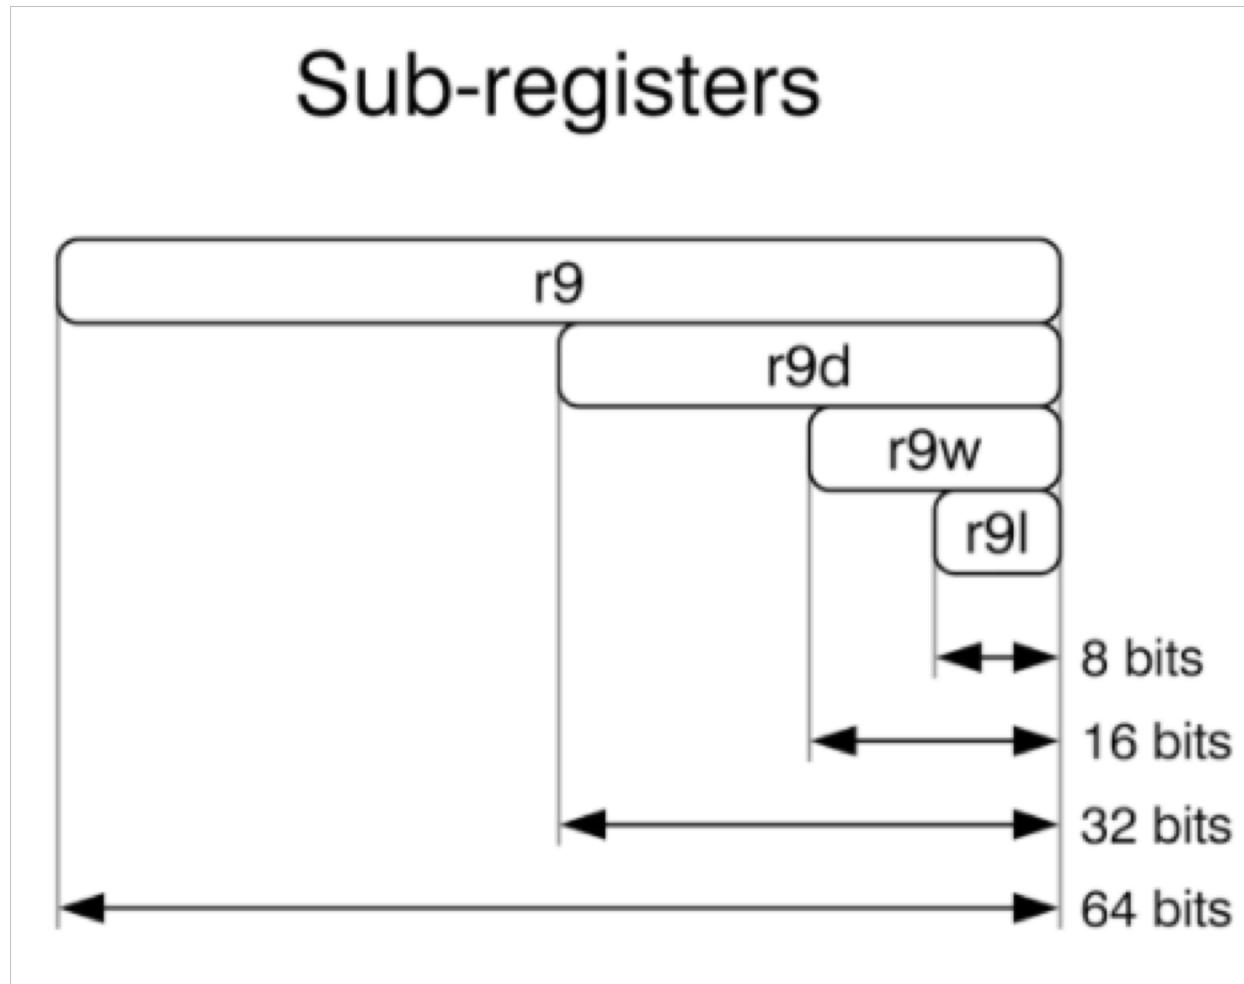
# L1 comparisons

- Saving the result of a comparison requires a few extra instructions

```
rdi <- rax <= rbx      ⇒  cmpq %rbx, %rax  
                           setle %dil  
                           movzbq %dil, %rdi
```

- cmpq updates a condition code in some hidden place (flags register)
- Then, we need to use setle to extract the condition code from this hidden place
- setle, however, needs an 8 bit register as its destination

# Intel sub-registers



# L1 comparisons

- Saving the result of a comparison requires a few extra instructions

```
rdi <- rax <= rbx      ⇒  cmpq %rbx, %rax  
                                setle %dil  
                                movzbq %dil, %rdi
```

- cmpq updates a condition code in some hidden place (flags register)
- Then, we need to use setle to extract the condition code from this hidden place
- setle, however, needs an 8 bit register as its destination
- So we use %dil here because that's an 8 bit register that overlaps with the lowest 8 bits of %rdi
- setle updates only those 8 bits; therefore we need movzbq to zero out the rest

# L1 comparisons

- Mapping register names to their 8-bit variants

**r10** → **r10b**

**r11** → **r11b**

**r12** → **r12b**

**r13** → **r13b**

**r14** → **r14b**

**r15** → **r15b**

**r8** → **r8b**

**r9** → **r9b**

**rax** → **al**

**rbp** → **bpl**

**rbx** → **bl**

**rcx** → **cl**

**rdi** → **dil**

**rdx** → **dl**

**rsi** → **sil**

# L1 comparisons

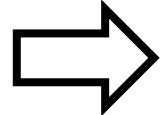
- Saving the result of a comparison requires a few extra instructions

```
rdi <- rax <= rbx      ⇒  cmpq %rbx, %rax  
                           setle %dil  
                           movzbq %dil, %rdi
```

- if we had < we'd need to use setg or setl (for less than or greater than)
- If we had = then we would use sete

# L1 comparisons with a constant

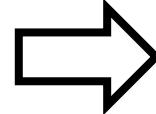
rdi <- rax <= 10



cmpq \$10, %rax  
setle %dil  
movzbq %dil, %rdi

# L1 comparisons with a constant

rdi <- 10 <= rex



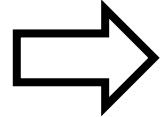
Must be a register

cmpq %rax, \$10  
setle %dil  
movzbq %dil, %rdi

Your compiler must handle this x86\_64-specific constraint

# L1 comparisons with a constant

rdi <- 10 <= rex



cmpq 10, %rax  
setge %dil  
movzbq %dil, %rdi

# L1 comparisons

So when we don't have any registers at all, we need to compute the answer at compile time and just use that

**rdi** <- 10 <= 11       $\Rightarrow$     **movq \$1, %rdi**

**rax** <- 12 <= 11       $\Rightarrow$     **movq \$0, %rax**

# L1 shifting operations

The shifting, **sop=**, operations also use the 8-bit registers, this time for their sources

**rdi <= rcx**       $\Rightarrow$     **salq %cl, %rdi**

**rdi >= 3**             $\Rightarrow$     **sarq \$3, %rdi**

The **l** is for “left shift” and the **r** stands for “right shift”.

# Labels and direct jumps

Labels and gotos are what you might guess; just replace the leading colon with an underscore and add a colon suffix when you define the label

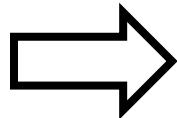
`:a_label` ⇒ `_a_label:`

`goto :a_label` ⇒ `jmp _a_label`

## Labels (2)

- When a label is stored in a memory location, you need to add “\$” before the label

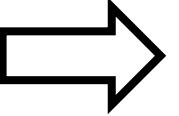
mem rsp -8 <- :myLabel



movq \$\_myLabel, -8(%rsp)

# Conditional jumps

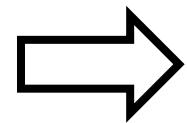
- We have the three same cases as for conditional comparisons
- Here, however, we use two jumps instead of storing the result in a register

cjump rax <= rdi :yes :no        cmpq %rdi, %rax  
    jle \_yes  
    jmp \_no

- For  $\leq$ , use jge (jump greater than or equal) or jle
- For  $<$ , use jg (jump greater than) or jl (jump less than)
- For  $=$ , use je

# Conditional jumps with constants

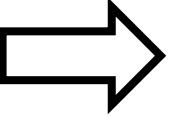
cjump 3 <= 1 :true :false



jmp \_false

# Conditional jumps (2)

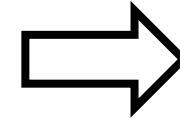
- We have the three same cases as for the previous conditional jump
- Here, however, we use only one jump

cjump rax <= rdi :yes            cmpq %rdi, %rax  
                                      jle \_yes

- For  $\leq$ , use jge (jump greater than or equal) or jle
- For  $<$ , use jg (jump greater than) or jl (jump less than)
- For  $=$ , use je

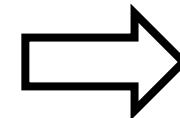
# Conditional jumps (2) with constants

cjump 1 <= 3 :true



jmp\_true

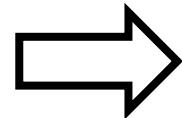
cjump 3 <= 1 :true



# The missing L1 CISC instruction

- The next instruction computes  $\text{rdi} + \text{rsi} * 4$

`rax @ rdi rsi 4`



`lea (%rdi, %rsi, 4), %rax`

# L1 instructions that modify rsp

- Function prologue (entry to a function)
- call and return instructions

# L1 function prologue

- The function prologue allocates locals
- For each stack variable: move the stack pointer by 8 bytes

```
(:myF  
 0 3  
  ...  
)
```



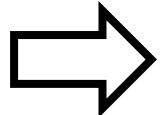
```
_myF:  
  subq $24, %rsp  #Allocate locals  
  ...
```

# L1 return instructions

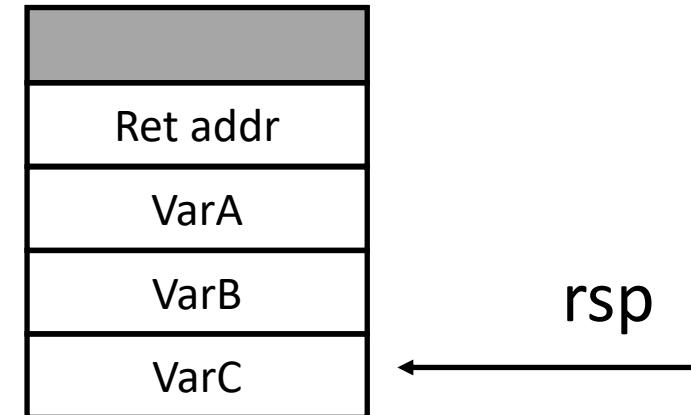
## The return instruction

- frees locals and ... (next slide)
- pops the return address from the stack and jumps to it

```
(:myF  
 0 3  
  ...  
  return  
)
```



```
...  
addq $24, %rsp  
retq
```

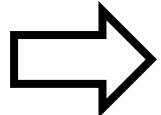


# L1 return instructions

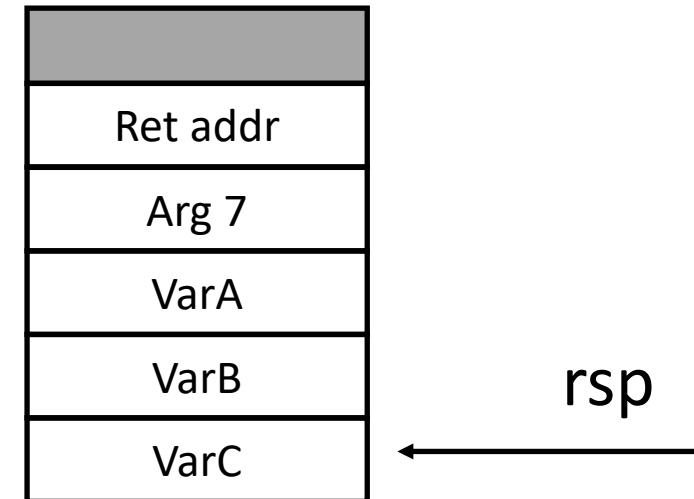
## The return instruction

- frees locals and stack arguments
- pops the return address from the stack and jumps to it

```
(:myF  
 7 3  
 ...  
 return  
)
```



```
...  
addq $32, %rsp  
retq
```



# L1 call instructions

Calls are translated differently depending on whether or not they invoke another L1 function

These calls are already considered differently in L1

- Calls to L1 functions: we have to store the return address

```
mem rsp -8 <- :f_ret  
call :myCallee  
:f_ret
```

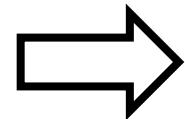
- Calls to the L1 runtime: we don't
- ```
call print 1
```

# L1 call instructions to L1 functions

The L1 call instructions to L1 functions

1. moves rsp based on the number of arguments  
and the return address
2. and then jumps to the callee

call :theCallee 11



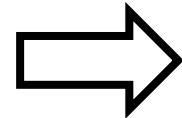
subq \$48, %rsp  
jmp \_theCallee

Why?  
 $(11 - 6) * 8 + 8$

# L1 indirect call instructions

- If call gets a register instead of a direct label, then the generated assembly code needs an extra asterisk

call rdi 0



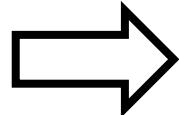
subq \$8, %rsp  
jmp \*%rdi

# L1 call instructions to runtime.c functions

The translation of these L1 call instructions

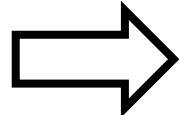
1. Does not need to change rsp
2. Relies on the Intel x86\_64 call instruction

call print 1



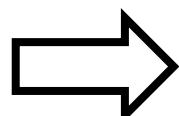
call print

call allocate 2



call allocate

call array-error 2



call array\_error

- It takes care of
1. identifying the return address
  2. storing the return address on the stack
  3. jumping to the callee

# Outline

- Setup
- From L1 to x86\_64
- Calling convention

# x86\_64 calling convention

- It is different than L1 calling convention

- Why does it matter for L1 programs?

call print 1

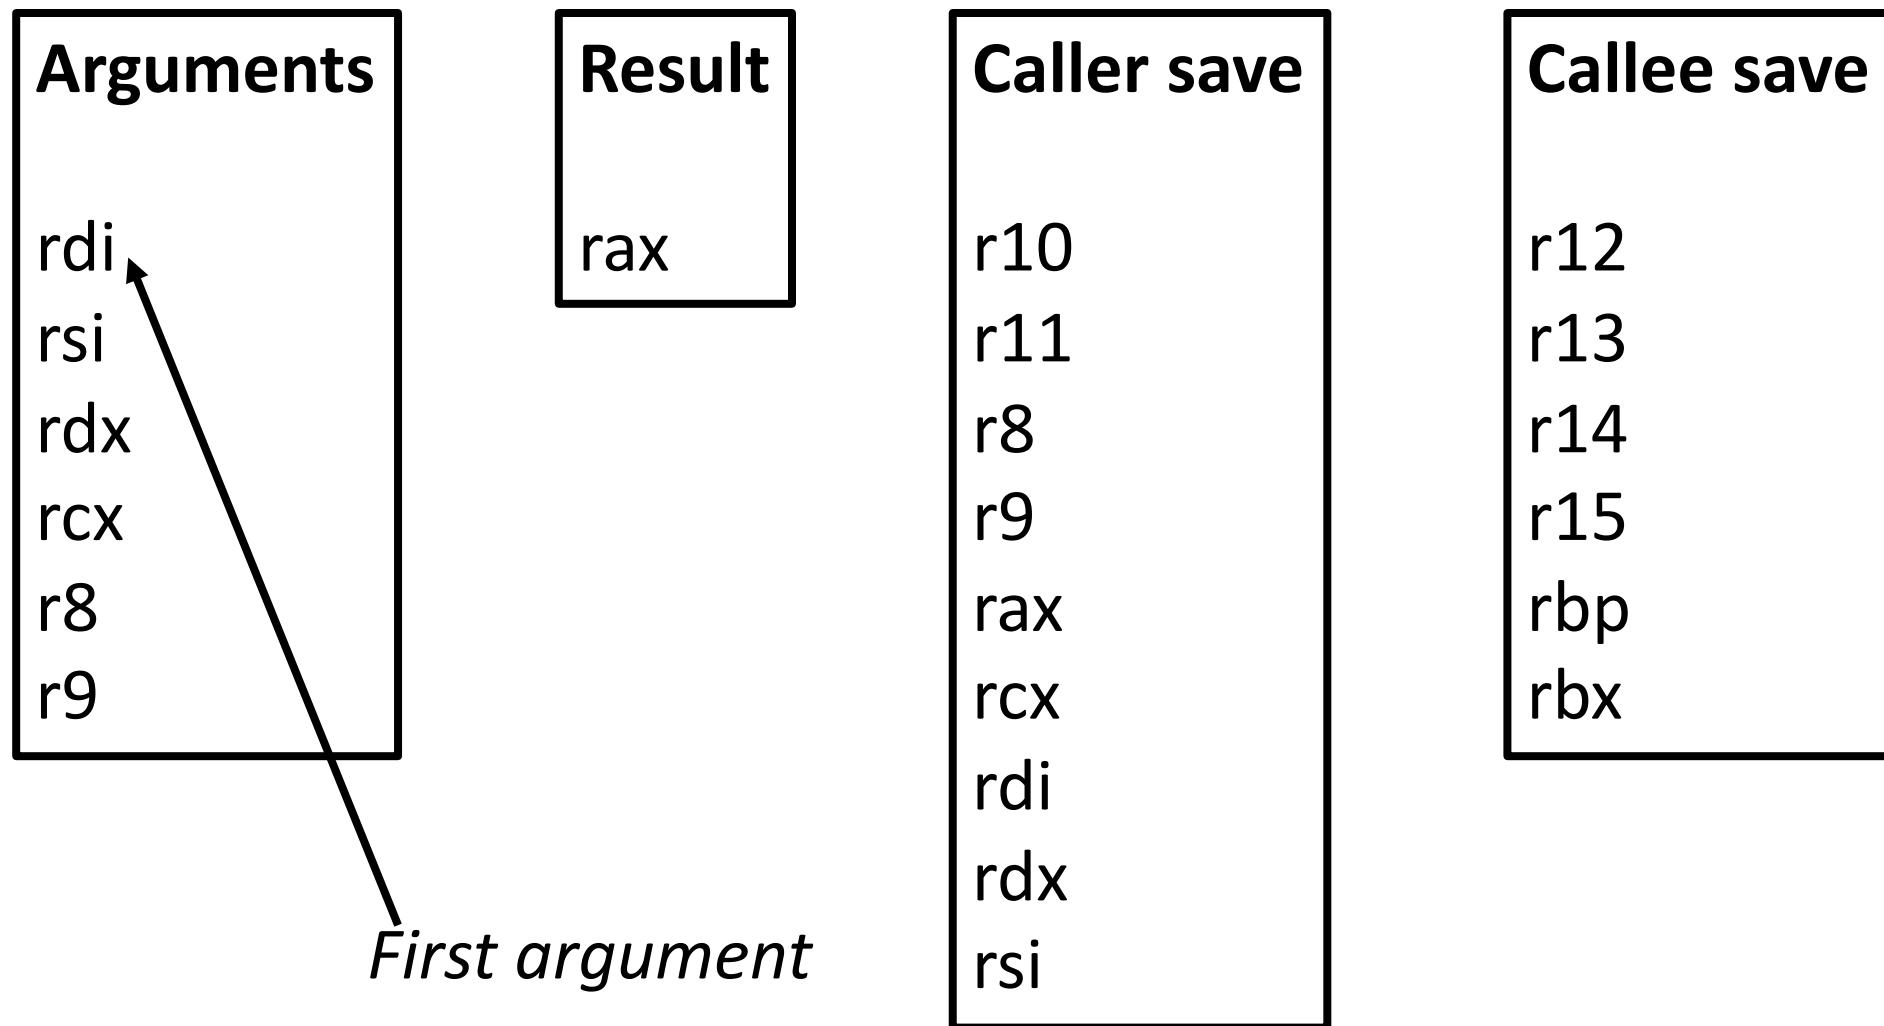
call allocate 2

call array\_error 2

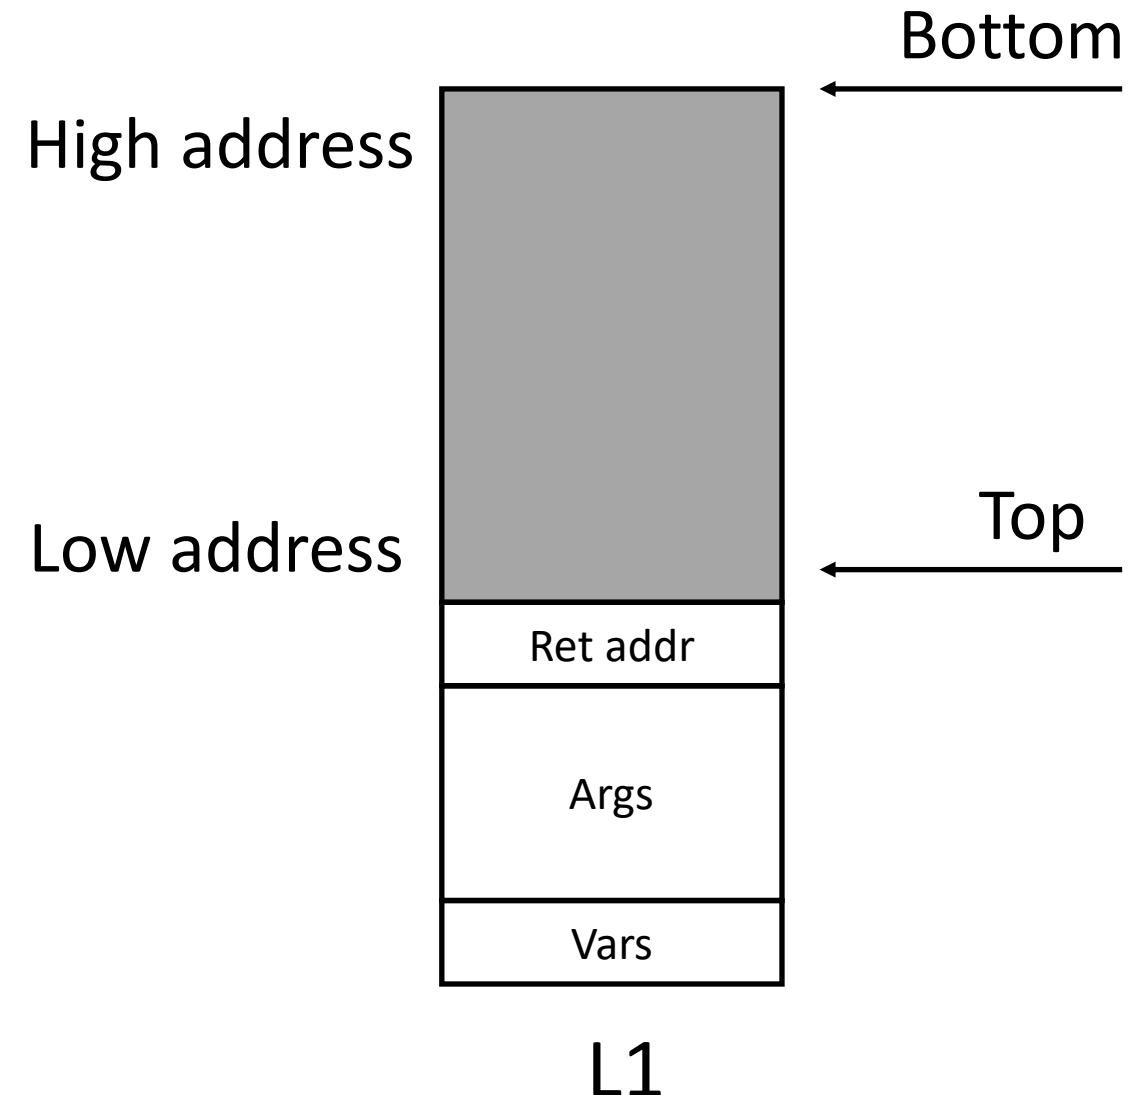
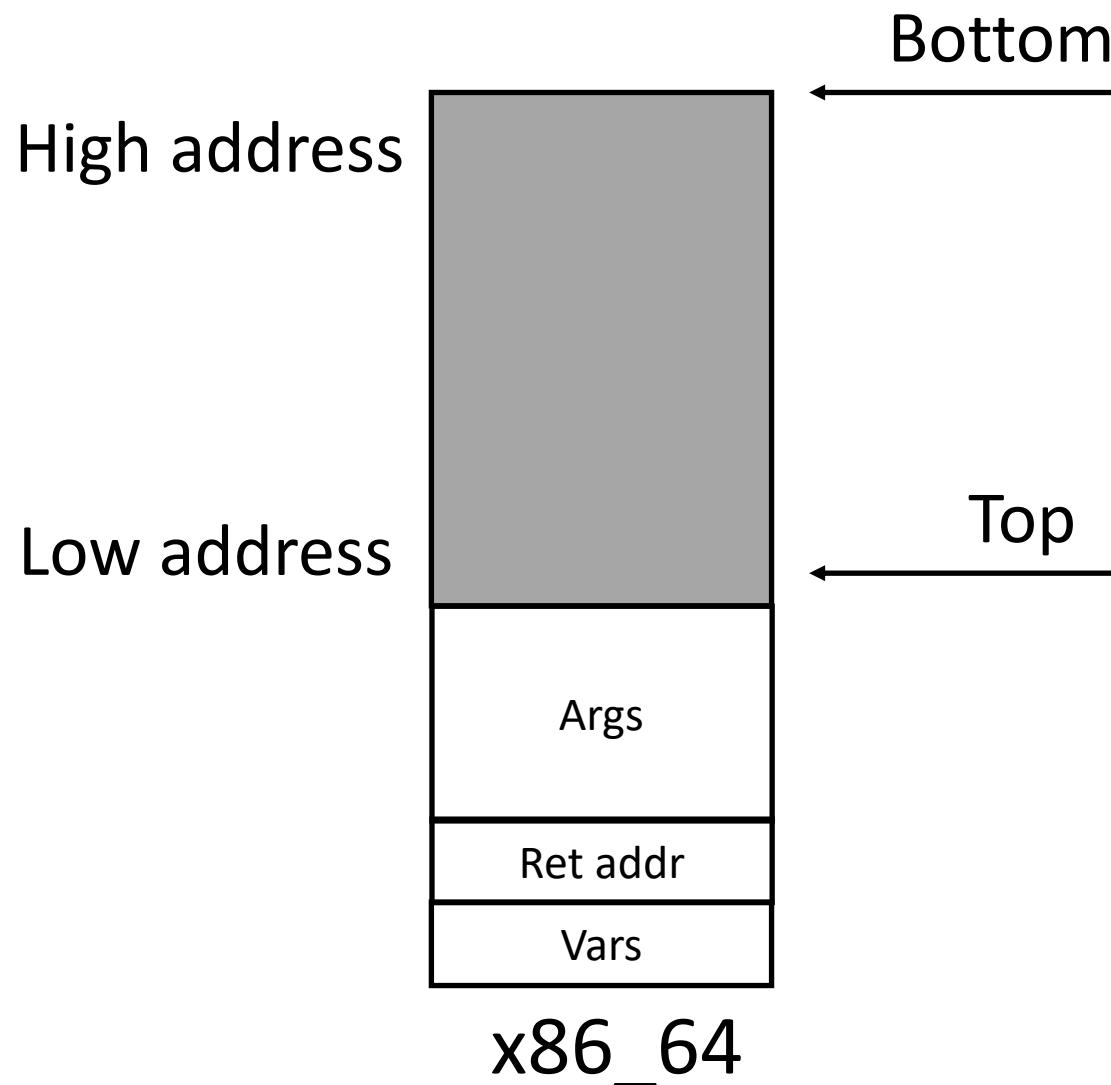
- runtime.c includes the body of these functions
- runtime.c is compiled with gcc,  
which follows x86\_64 calling convention

**Why does it work then?**

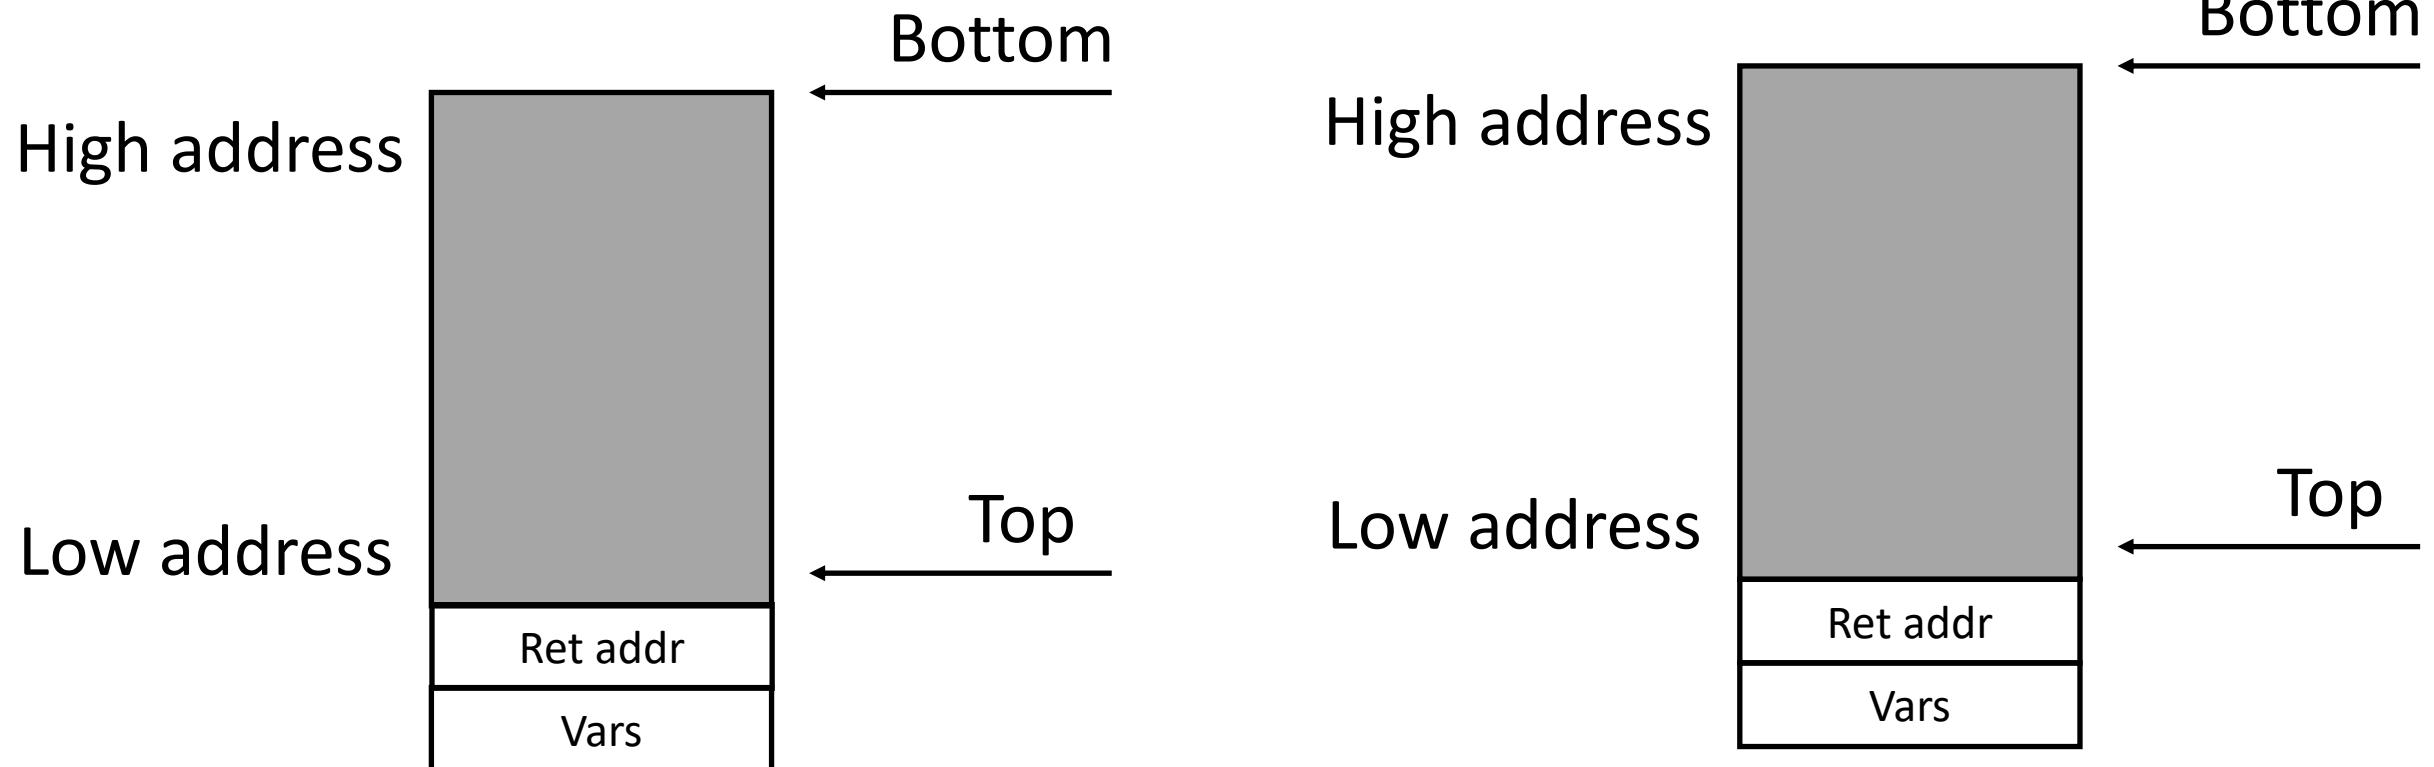
# Registers (same for L1)



# The stack (different compared to L1)

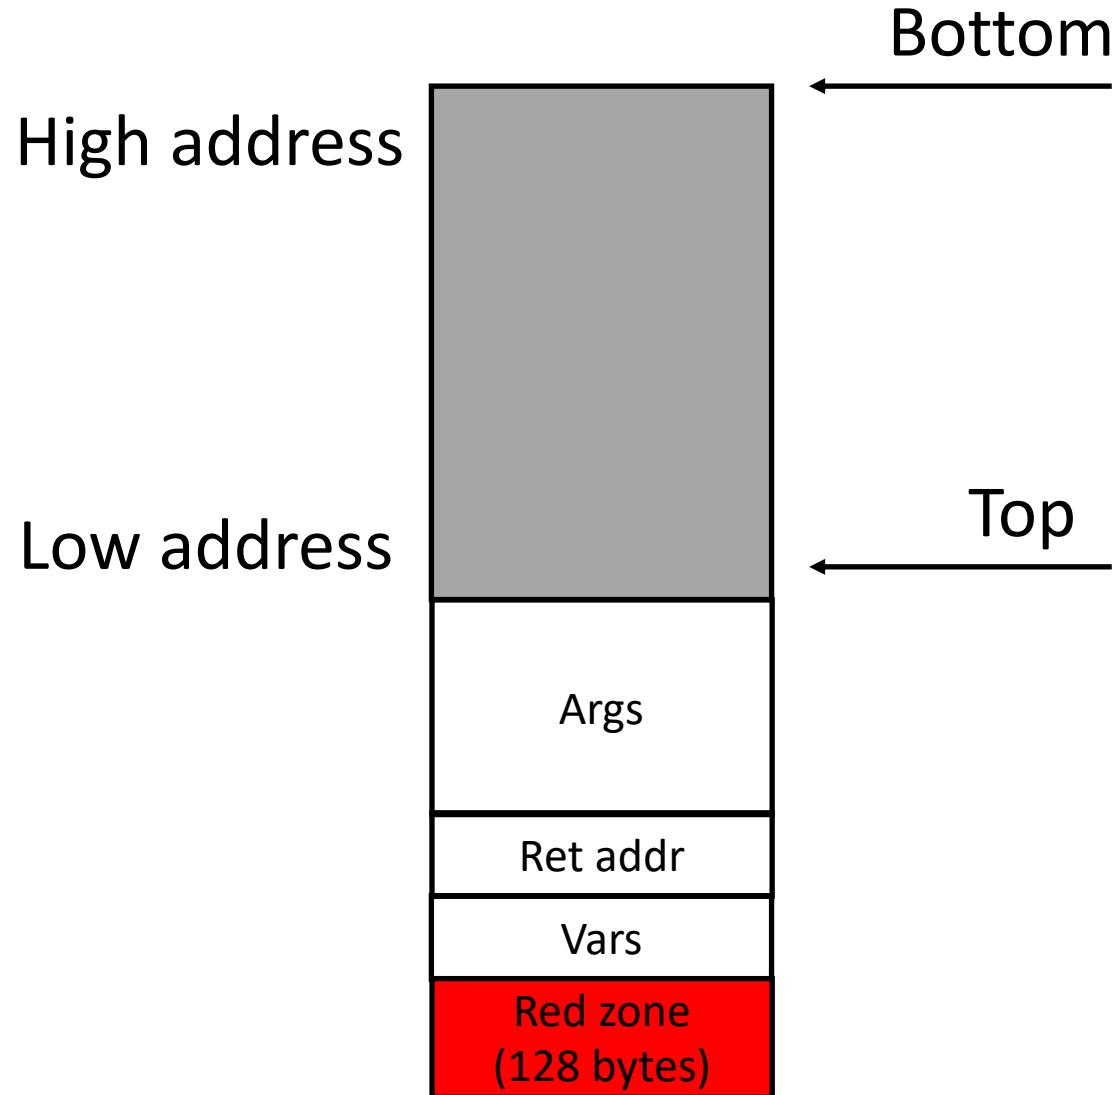


# The stack for runtime.c

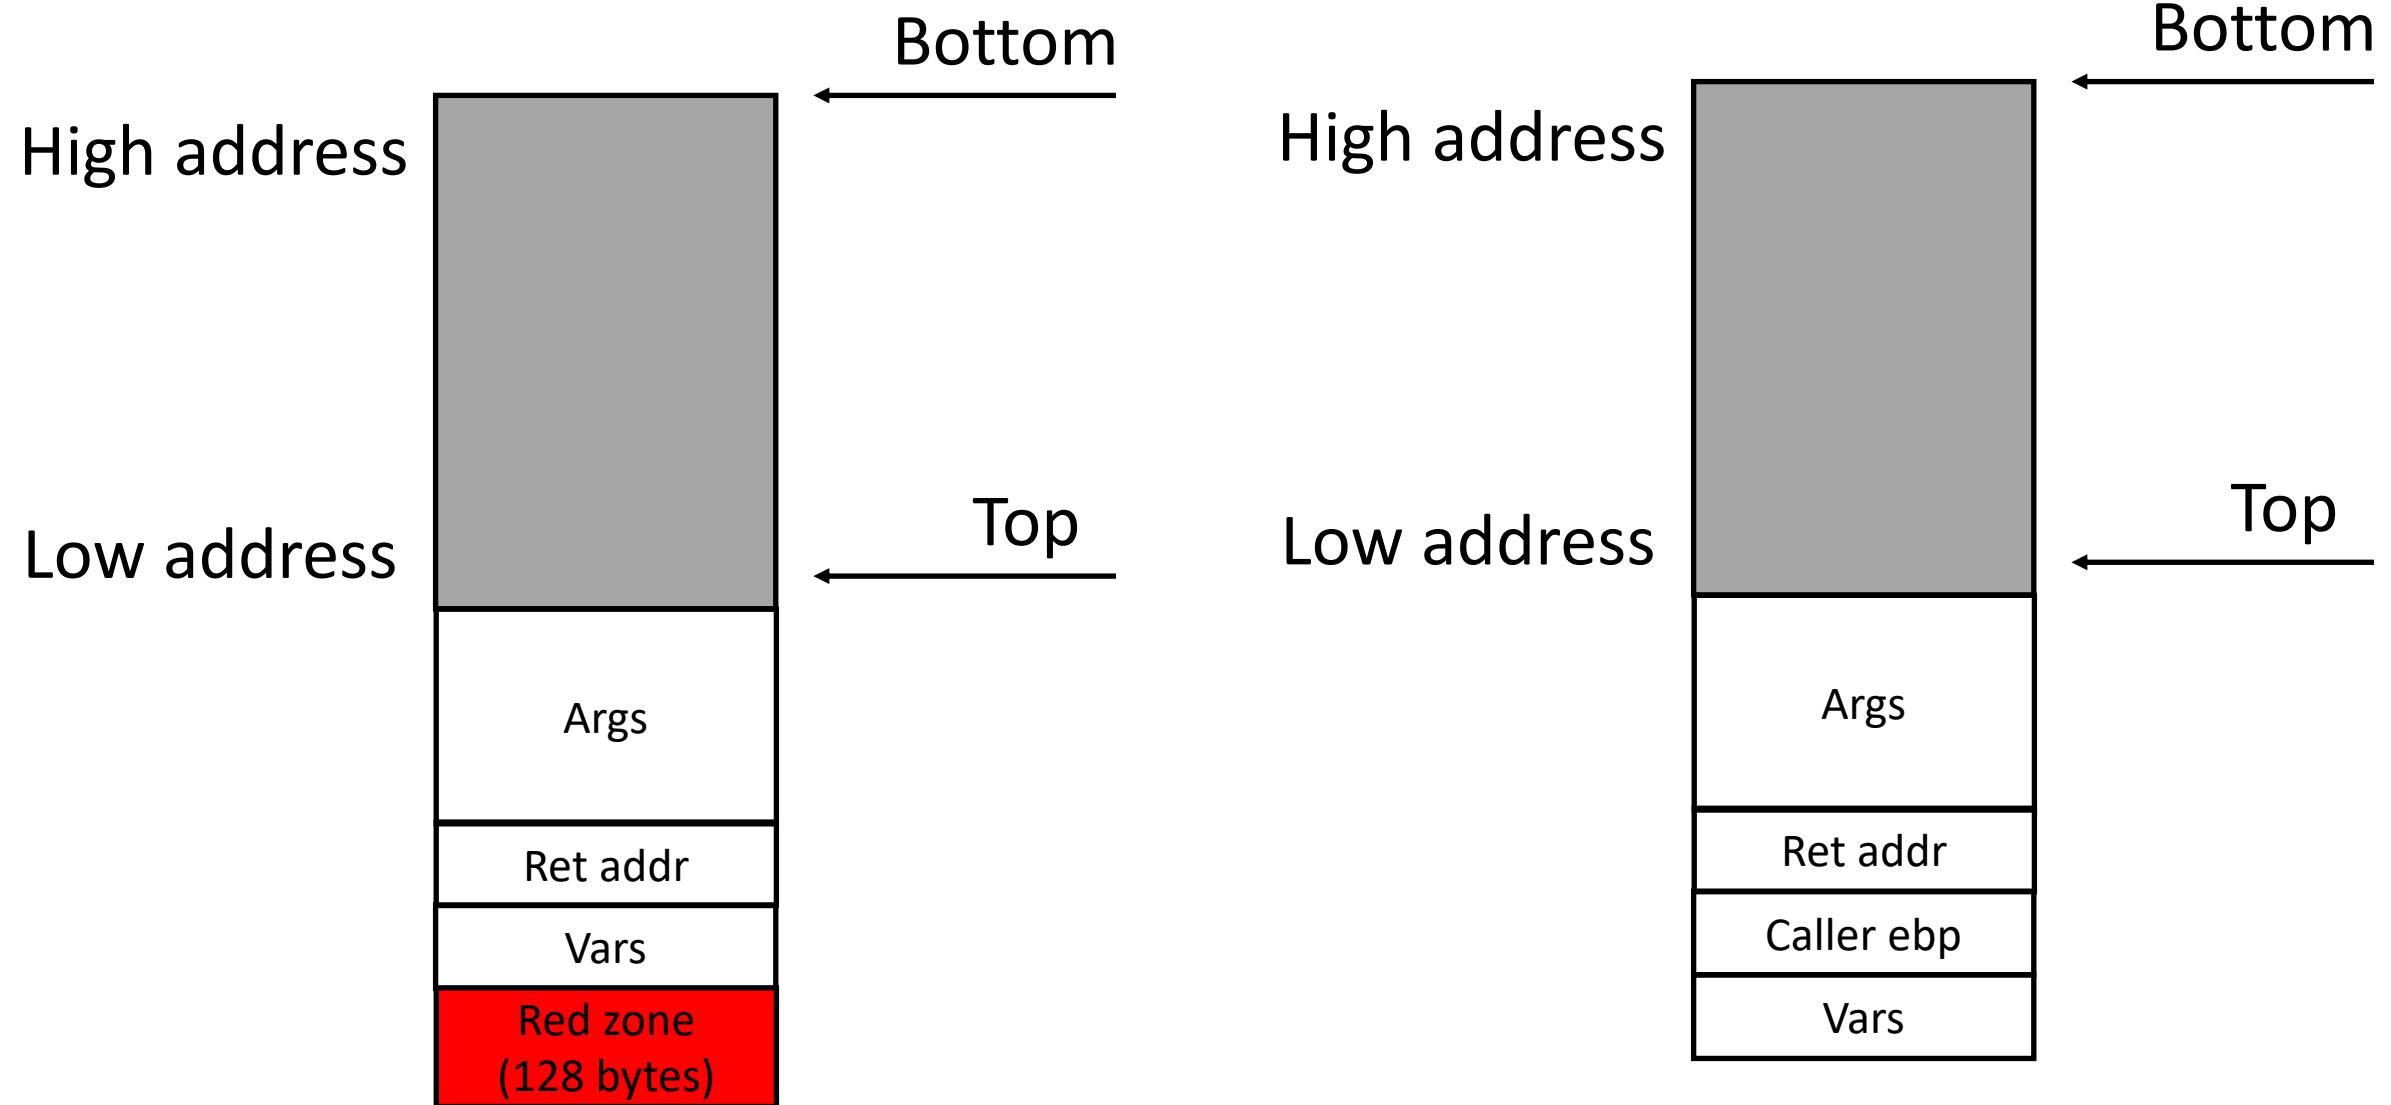


The callee is responsible for allocating and deallocating Vars

# More about x86\_64 calling convention



# x86\_64 vs. x86 calling convention



# Homework 0

- Develop the L1 compiler  
to translate L1 programs to x86\_64 binaries
  - You must follow the translation specified by these slides
  - You must be able to pass all tests  
`cd L1 ; make test`
- Deadline: see Canvas